

# Applied\_Stat\_2\_Lab2

```
library(opendatatoronto)
library(tidyverse)
library(stringr)
library(skimr) # EDA
library(visdat) # EDA
library(janitor)
library(lubridate)
library(ggrepel)
```

```
res <- list_package_resources("996cfe8d-fb35-40ce-b569-698d51fc683b") # obtained code from
res <- res |> mutate(year = str_extract(name, "202.?"))
delay_2022_ids <- res |> filter(year==2022) |> select(id) |> pull()
```

```
delay_2022 <- get_resource(delay_2022_ids)
```

```
# make the column names nicer to work with
delay_2022 <- clean_names(delay_2022)
```

```
delay_codes <- get_resource("3900e649-f31e-4b79-9f20-4731bbfd94f7")
```

New names:

```
* `` -> `...1`
* `CODE DESCRIPTION` -> `CODE DESCRIPTION...3`
* `` -> `...4`
* `` -> `...5`
* `CODE DESCRIPTION` -> `CODE DESCRIPTION...7`
```

```
delay_data_codebook <- get_resource("ca43ac3d-3940-4315-889b-a9375e7b8aa4")
```

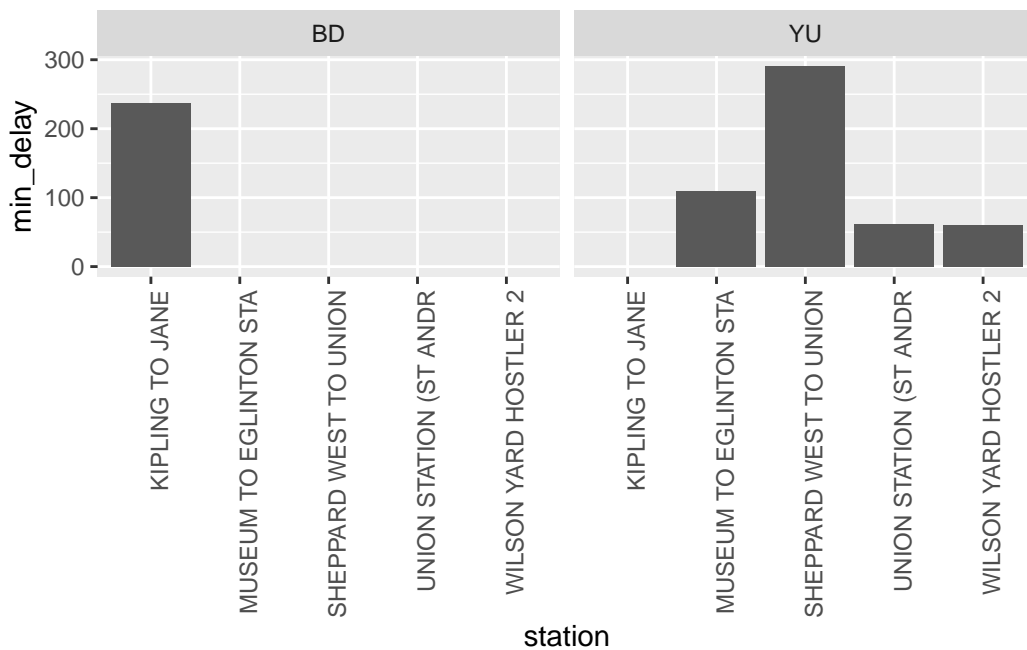
```
delay_2022 <- delay_2022 |>
  mutate(station_clean = ifelse(str_starts(station, "ST"), word(station, 1,2), word(station,
```

## Question 1

```
#Calculate mean delays and sort descending
mean_delays <- delay_2022 |> group_by(station) |> summarize(mean_delay = mean(min_delay),

#Take 5 highest mean delays and add the rest of the data
highest_delays <- head(mean_delays, 5)
delay_2022_top_stations <- delay_2022 %>%
  filter(station %in% highest_delays$station)

ggplot(delay_2022_top_stations, aes(x = station, y = min_delay)) +
  geom_bar(stat = "identity") +
  theme(axis.text.x = element_text(angle = 90, hjust =1)) +
  facet_wrap(~line)
```



## Question 2

```
delay_2022 <- delay_2022 |>
  left_join(delay_codes |> rename(code = `SUB RMENU CODE`, code_desc = `CODE DESCRIPTION.
```

Joining with `by = join\_by(code)`

```
#Filter data by top 50% of delays
delay_2022_top_0.5 <- delay_2022 |>
  filter(min_delay>0)|>
  group_by(code)|>
  summarise(no_rows = length(code))|>
  arrange(-no_rows)|>
  mutate(cumulative_sum = cumsum(no_rows))|>
  mutate(half_sum = sum(no_rows)/2)|>
  filter(cumulative_sum<=half_sum)

frequent_delay_codes <- delay_2022_top_0.5$code

lm_table_delay_code <- delay_2022 |>
  filter(min_delay>0 & (code %in% frequent_delay_codes))
#Linear model with line and code as covariates
delay_model <- lm(min_delay ~ line + code, data = lm_table_delay_code)
summary(delay_model)
```

Call:

```
lm(formula = min_delay ~ line + code, data = lm_table_delay_code)
```

Residuals:

Min	1Q	Median	3Q	Max
-10.475	-2.450	-1.072	0.890	227.525

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	5.7698	0.3485	16.554	< 2e-16 ***
lineSHP	1.3899	0.5828	2.385	0.017132 *
lineYU	-0.3203	0.2521	-1.270	0.204022
codeMUIR	1.5470	0.4432	3.491	0.000486 ***

codeMUPAA	-1.6602	0.3741	-4.438	9.3e-06	***
codePUOPO	-0.9396	0.3405	-2.759	0.005814	**
codeSUDP	0.9928	0.3344	2.969	0.003003	**
codeSUO	5.1117	0.4381	11.667	< 2e-16	***
codeSUUT	7.7057	0.4069	18.938	< 2e-16	***
codeTUNOA	-1.3775	0.3954	-3.484	0.000499	***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.38 on 4396 degrees of freedom

(1 observation deleted due to missingness)

Multiple R-squared: 0.1668, Adjusted R-squared: 0.1651

F-statistic: 97.8 on 9 and 4396 DF, p-value: < 2.2e-16

Our model suggests that the line “SHP” contributes significantly to delay minutes, which was not seen at all in the Explanatory Data Analysis in Question 1. This is because in Question 1 we found the five stations with highest mean delay. Naturally, the highest delay times were produced by outlier accidents that caused significantly longer delays than the total average is. Hence the stations or lines with the most frequent albeit shorter delays were not captured in the EDA. This is also explained by the fact that the most frequent delays do not contribute to the highest delays, which is also suggested by the relatively low beta coefficient estimates in our model.

### Question 3

```
#Data preprocessing
all_data <- search_packages("campaign")
campaign_id <- all_data$id
resource <- list_package_resources(campaign_id[1])
campaign_data <- get_resource('8b42906f-c894-4e93-a98e-acac200f34a4')
```

New names:

New names:

New names:

New names:

New names:

New names:

New names:

\* `` -> `...2`

\* `` -> `...3`

```
campaign_data_mayoral <- campaign_data[[2]]
colnames(campaign_data_mayoral) <- as.character(campaign_data_mayoral[1,])
campaign_data_mayoral <- campaign_data_mayoral[-1,]
rownames(campaign_data_mayoral) <- NULL
campaign_data_mayoral <- clean_names(campaign_data_mayoral)
```

## Question 4

We skim through the data using the skim function.

```
skim(campaign_data_mayoral)
```

Table 1: Data summary

Name	campaign_data_mayoral
Number of rows	10199
Number of columns	13
Column type frequency:	
character	13
Group variables	None

### Variable type: character

skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
contributors_name	0	1	4	31	0	7545	0
contributors_address	10197	0	24	26	0	2	0
contributors_postal_code	0	1	7	7	0	5284	0
contribution_amount	0	1	1	18	0	209	0
contribution_type_desc	0	1	8	14	0	2	0
goods_or_service_desc	10188	0	11	40	0	9	0
contributor_type_desc	0	1	10	11	0	2	0
relationship_to_candidate	10166	0	6	9	0	2	0
president_business_manager	10197	0	13	16	0	2	0
authorized_representative	10197	0	13	16	0	2	0
candidate	0	1	9	18	0	27	0
office	0	1	5	5	0	1	0
ward	10199	0	NA	NA	0	0	0

---

skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
---------------	-----------	---------------	-----	-----	-------	----------	------------

---

There are many blank columns or columns with missing values such as ‘contributors\_address’, ‘authorized\_representative’, ‘president\_business\_manager’ and so on. Furthermore, there are a couple of variables that should be factors and some that should be numerical.

```
not_all_na <- function(x) all(!is.na(x))
campaign_data_mayoral <- campaign_data_mayoral|>
  select(where(not_all_na))
campaign_data_mayoral
```

```
# A tibble: 10,199 x 7
  contributors_name contributors_postal_code contribution_amount
  <chr>             <chr>                                <chr>
1 A D'Angelo, Tullio M6A 1P5                                300
2 A Strazar, Martin M2M 3B8                                300
3 A'Court, K Susan M4M 2J8                                  36
4 A'Court, K Susan M4M 2J8                                  100
5 A'Court, K Susan M4M 2J8                                  100
6 Aaron, Robert B M6B 1H7                                    250
7 Abadi, Babak M5S 2W7                                       500
8 Abadi, Babak M5S 2W7                                       500
9 Abadi, David M5S 2W7                                       300
10 Abate, Frank L4H 2K7                                       150
# i 10,189 more rows
# i 4 more variables: contribution_type_desc <chr>,
# contributor_type_desc <chr>, candidate <chr>, office <chr>
```

```
campaign_data_mayoral$contributor_type_desc <- as.factor(campaign_data_mayoral$contributor_type_desc)
campaign_data_mayoral$contribution_type_desc <- as.factor(campaign_data_mayoral$contribution_type_desc)
campaign_data_mayoral$contribution_amount <- as.numeric(campaign_data_mayoral$contribution_amount)
campaign_data_mayoral
```

```
# A tibble: 10,199 x 7
  contributors_name contributors_postal_code contribution_amount
  <chr>             <chr>                                <dbl>
1 A D'Angelo, Tullio M6A 1P5                                300
2 A Strazar, Martin M2M 3B8                                300
```

```

3 A'Court, K Susan      M4M 2J8      36
4 A'Court, K Susan      M4M 2J8     100
5 A'Court, K Susan      M4M 2J8     100
6 Aaron, Robert B       M6B 1H7     250
7 Abadi, Babak          M5S 2W7     500
8 Abadi, Babak          M5S 2W7     500
9 Abadi, David          M5S 2W7     300
10 Abate, Frank          L4H 2K7     150
# i 10,189 more rows
# i 4 more variables: contribution_type_desc <fct>,
#   contributor_type_desc <fct>, candidate <chr>, office <chr>

```

## Question 5

```
campaign_data_mayoral |> arrange(-contribution_amount)
```

```

# A tibble: 10,199 x 7
  contributors_name contributors_postal_code contribution_amount
  <chr>             <chr>                                <dbl>
1 Ford, Doug        M9A 2C3                                508225.
2 Ford, Rob         M9A 3G9                                78805.
3 Ford, Doug        M9A 2C3                                50000
4 Ford, Rob         M9A 3G9                                50000
5 Ford, Rob         M9A 3G9                                50000
6 Goldkind, Ari     M5P 1P5                                23624.
7 Ford, Rob         M9A 3G9                                20000
8 Ford, Rob         M9A 3G9                                12210
9 Di Paola, Rocco   M3H 2T1                                 6000
10 Thomson, Sarah   M4W 2X6                                4426.
# i 10,189 more rows
# i 4 more variables: contribution_type_desc <fct>,
#   contributor_type_desc <fct>, candidate <chr>, office <chr>


```

```

ggplot(data = campaign_data_mayoral, aes(x=contribution_amount))+
  geom_dotplot()

```


Bin width defaults to 1/30 of the range of the data. Pick better value with ``binwidth``.



```
campaign_data_mayoral_contribution_distribution <-  
  campaign_data_mayoral |> filter(contribution_amount < 100000)  
  
ggplot(data=campaign_data_mayoral_contribution_distribution, aes(x=contribution_amount))+  
  geom_dotplot()
```

Bin width defaults to 1/30 of the range of the data. Pick better value with  
`binwidth`.






```
campaign_data_mayoral_contribution_distribution_2 <-
  campaign_data_mayoral |> filter(contribution_amount < 10000)


ggplot(data=campaign_data_mayoral_contribution_distribution_2, aes(x=contribution_amount))
  geom_histogram()
```

`stat\_bin()` using `bins = 30`. Pick better value with `binwidth`.



```
ggplot(data = campaign_data_mayoral, aes(x = candidate,  
                                           y = contribution_amount,  
                                           color = contribution_type_desc ))+  
  geom_point()+  
  geom_smooth()+  
  theme(axis.text.x = element_text(angle = 90, hjust = 1))
```

`geom\_smooth()` using method = 'gam' and formula = 'y ~ s(x, bs = "cs")'



Looking at the plots above, the following outliers can be found: The biggest outlier is the donation of 500,000\$ by Doug Ford. Then there are also a few other donations close to 100,000\$ done by candidates Ryan Emond and Rob Ford. After filtering for donations less than 100,000\$ we see a few outlier donations of above 20,000\$. Thus we filter once more for donations less than 10,000\$ to get a better sense of the majority of the data.

## Question 6

```
candidate_contri <- campaign_data_mayoral|>
  group_by(candidate) |>
  summarise(
    total_contri = sum(contribution_amount, na.rm = TRUE),
    mean_contri = mean(contribution_amount, na.rm = TRUE),
    contri_count = n()
  )

top_total_contri <- candidate_contri |>
  arrange(-total_contri)|>
  select(candidate,total_contri)|>
  head(5)

top_mean_contri <- candidate_contri |>
```

```

      arrange(-mean_contri)|>
      select(candidate,mean_contri)|>
      head(5)
top_contri_count <- candidate_contri |>
      arrange(-contri_count)|>
      select(candidate,contri_count)|>
      head(5)

```

```
top_total_contri
```

```

# A tibble: 5 x 2
  candidate      total_contri
  <chr>          <dbl>
1 Tory, John    2767869.
2 Chow, Olivia  1638266.
3 Ford, Doug    889897.
4 Ford, Rob     387648.
5 Stintz, Karen 242805

```

```
top_mean_contri
```

```

# A tibble: 5 x 2
  candidate      mean_contri
  <chr>          <dbl>
1 Sniedzins, Erwin    2025
2 Syed, Himy          2018
3 Ritch, Carlisle    1887.
4 Ford, Doug         1456.
5 Clarke, Kevin       1200

```

```
top_contri_count
```

```

# A tibble: 5 x 2
  candidate      contri_count
  <chr>          <int>
1 Chow, Olivia    5708
2 Tory, John      2602
3 Ford, Doug       611
4 Ford, Rob        538
5 Soknacki, David   314

```

## Question 7

```
non_candidate_contri <- campaign_data_mayoral |>
  filter(contributors_name != candidate)

non_candidate_contri <- non_candidate_contri|>
  group_by(candidate) |>
  summarise(
    total_contri_popular = sum(contribution_amount, na.rm = TRUE),
    mean_contri_popular = mean(contribution_amount, na.rm = TRUE),
    contri_count_popular = n()
  )

top_total_contri_popular <- non_candidate_contri |>
  arrange(-total_contri_popular)|>
  select(candidate,total_contri_popular)|>
  head(5)

top_mean_contri_popular <- non_candidate_contri |>
  arrange(-mean_contri_popular)|>
  select(candidate,mean_contri_popular)|>
  head(5)

top_contri_count_popular <- non_candidate_contri |>
  arrange(-contri_count_popular)|>
  select(candidate,contri_count_popular)|>
  head(5)

top_total_contri_popular

# A tibble: 5 x 2
  candidate      total_contri_popular
  <chr>          <dbl>
1 Tory, John    2765369.
2 Chow, Olivia  1634766.
3 Ford, Doug    331173.
4 Stintz, Karen 242805
5 Ford, Rob     174510.

top_mean_contri_popular

# A tibble: 5 x 2
  candidate      mean_contri_popular
```

	<chr>	<dbl>
1	Ritch, Carlie	1887.
2	Sniedzins, Erwin	1867.
3	Tory, John	1063.
4	Gardner, Norman	1000
5	Tiwari, Ramnarine	1000

```
top_contri_count_popular
```

```
# A tibble: 5 x 2
  candidate      contri_count_popular
  <chr>          <int>
1 Chow, Olivia    5706
2 Tory, John      2601
3 Ford, Doug       608
4 Ford, Rob        531
5 Soknacki, David  314
```

## Question 8

```
multiple_contri <- campaign_data_mayoral |>
  group_by(contributors_name) |>
  summarise(unique_candidates = n_distinct(candidate))

multiple_contri_count <- sum(multiple_contri$unique_candidates > 1)
multiple_contri_count
```

```
[1] 184
```